

Candidate review with explicit tie semantics

Aadi Ajeesh Nair and Ryan Gomez

Constructed arithmetic evidence. No real-sky performance or discovery claim.

The question

Can changing CSV row order alter a score when equal priorities cross a review-budget boundary?

The declared policy

Keep higher scores. For tied scores, use expected positives under uniform selection. Ranking cannot replace catalogue evidence or human review.


What was checked

24 row permutations x two score columns x four budgets. All 192 outputs match exact rational precision and recall.

What this supports

A reproducible software demonstration, explicit metric semantics and a reviewable evidence trail.

One tied cohort across four review budgets



Four constructed objects, two positive labels, all incumbent scores equal. A selected single row can be positive or negative; the declared expectation remains one half.

What remains open

Dated photometry, independent labels and prospective validation are still needed for sky completeness, purity, discovery yield or comparative model benefit.